

Classical and Modern Approaches for Solving the Steiner Tree Problem

Teodora Dobos

Technical University of Munich

June 21, 2022

Steiner Tree Problem

Given:

- ▶ $G = (V, E)$ undirected
- ▶ cost function $c : E \rightarrow \mathbb{Q}^+$
- ▶ terminals set $R \subseteq V$

Objective:

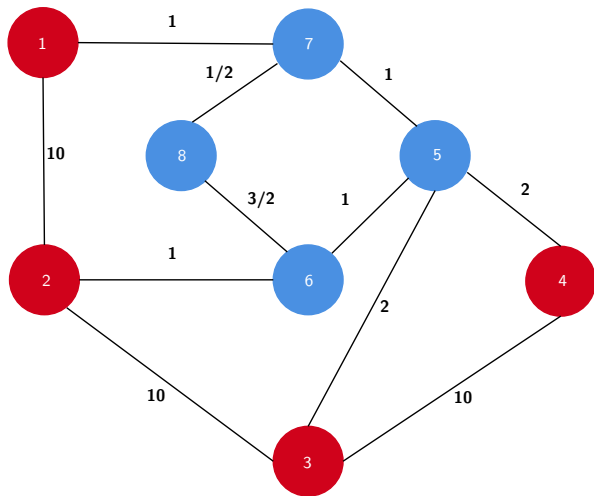
- ▶ find tree that connects all terminals and has minimum weight

1. Repetitive Shortest Path Heuristic
2. Kou-Markowsky-Berman Algorithm (Metric Closure)
3. Mehlhorn Algorithm
4. Primal-Dual Method

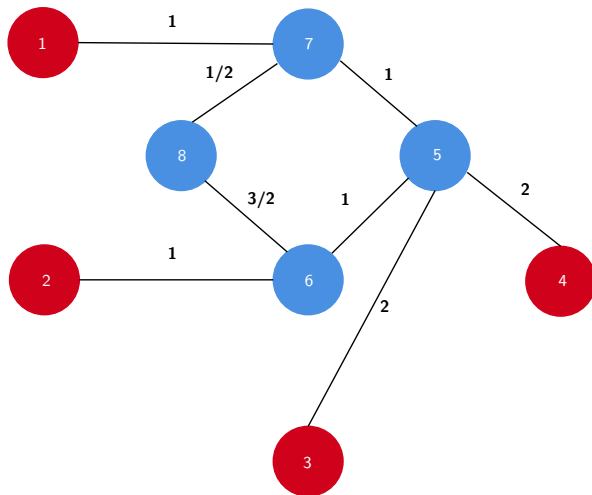
1. Repetitive Shortest Path Heuristic → implemented ✓
 2. Kou-Markowsky-Berman Algorithm → NetworkX package
 3. Mehlhorn Algorithm → implemented ✓
 4. Primal-Dual Method → implemented ✓
- $2 \left(1 - \frac{1}{|R|}\right)$ approximation algorithms

Repetitive Shortest Path Heuristic

Input graph



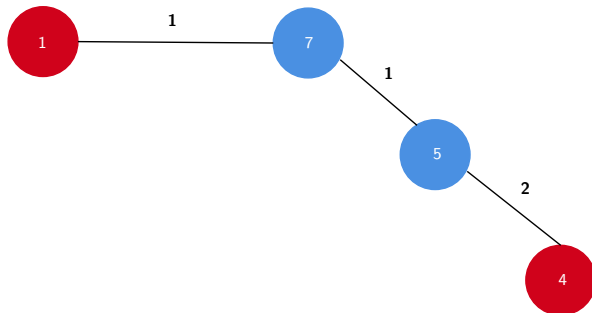
Compute shortest paths with sources $r \in R$



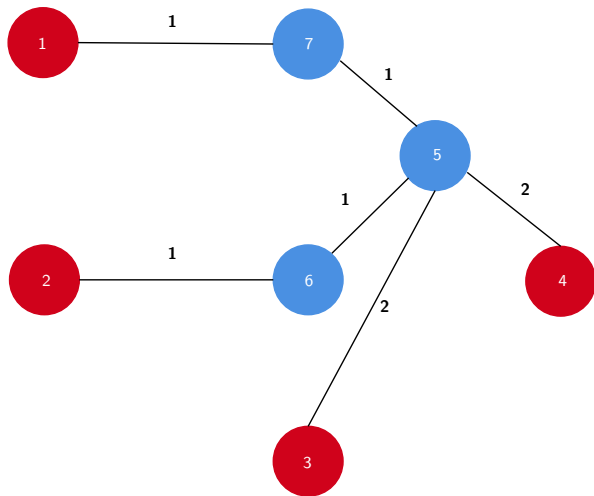
Choose a random terminal: $T = (\{r\}, \emptyset)$



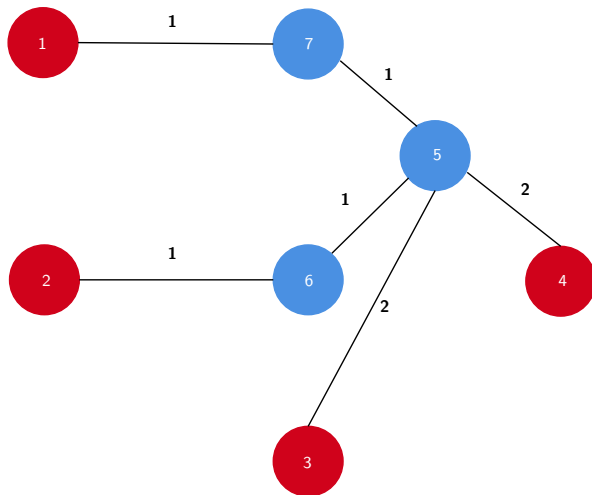
Choose closest $u \in R$ and connect it to the component via its shortest path



Repeat until T contains all terminals



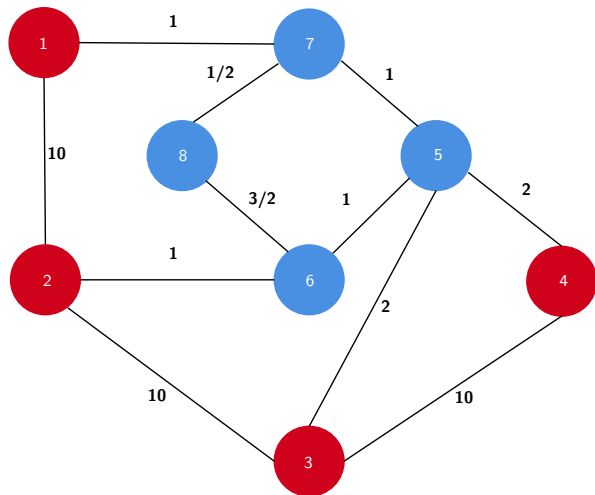
Remove all non-terminals with degree 1 (no change in this example)



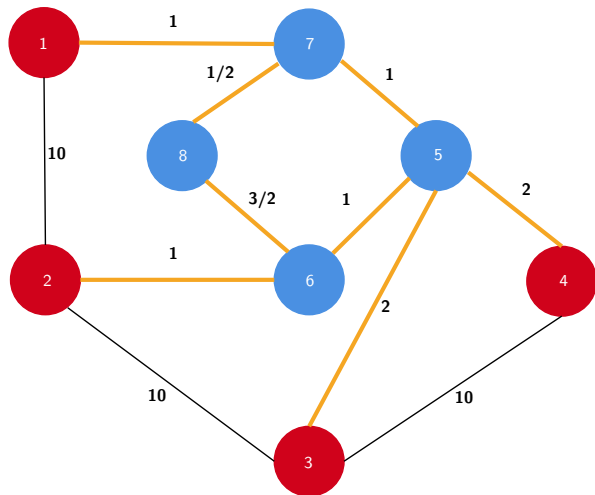
1. compute shortest paths with sources $r \in R$
2. $T = (\{r\}, \emptyset)$ with arbitrary $r \in R$
3. grow T to obtain a feasible solution
 - 3.1 choose closest $u \in R$
 - 3.2 add minimum path $u \rightarrow T$
4. remove all non-terminals with degree 1

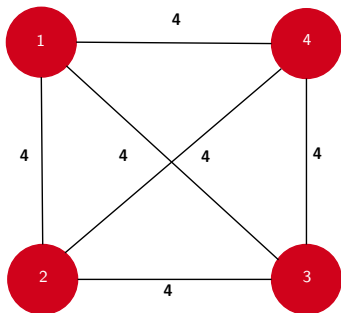
Kou-Markowsky-Berman Algorithm (Metric Closure)

Input graph G

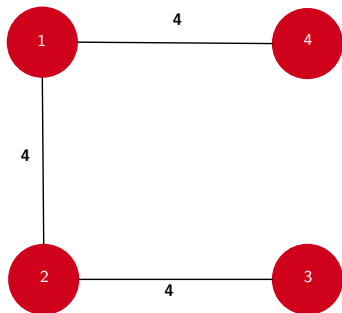


Compute the metric closure of G

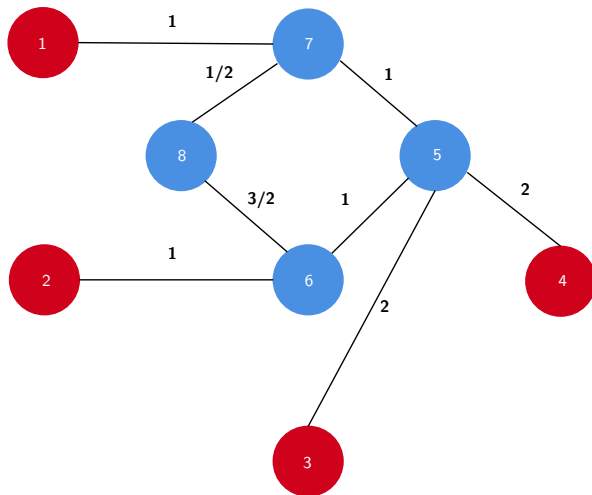




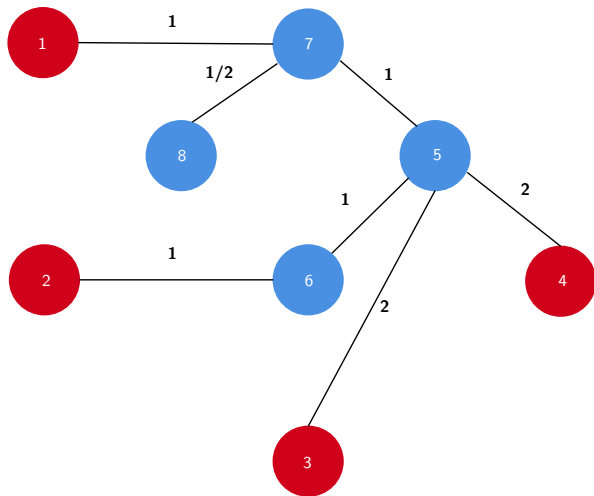
Compute a minimum spanning tree $G_2 = MST(G_1)$



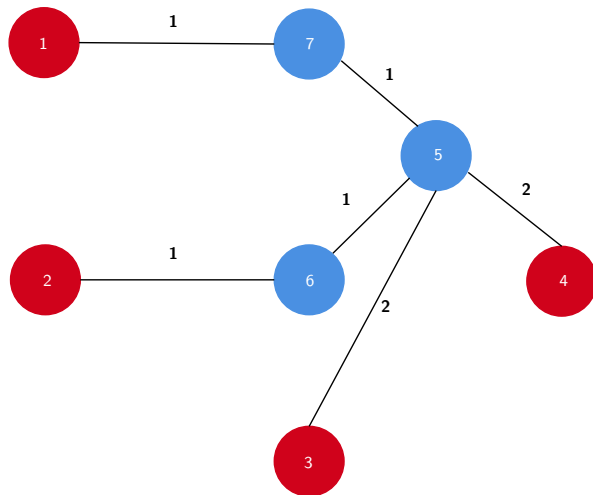
Replace edges in G_2 with corresp. shortest paths in $G \rightarrow G_3$



Compute a minimum spanning tree $G_4 = MST(G_3)$



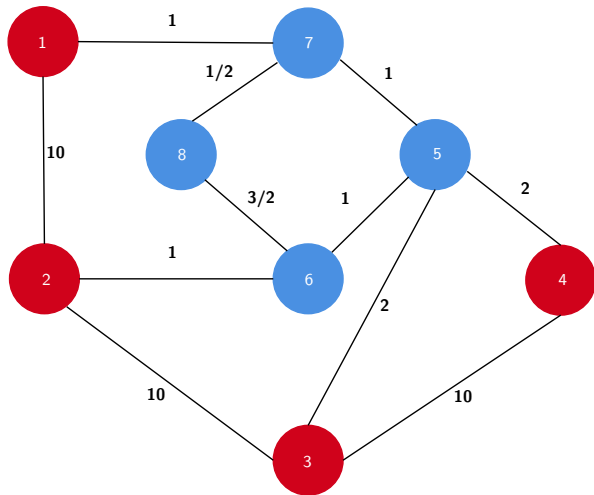
Delete unnecessary edges



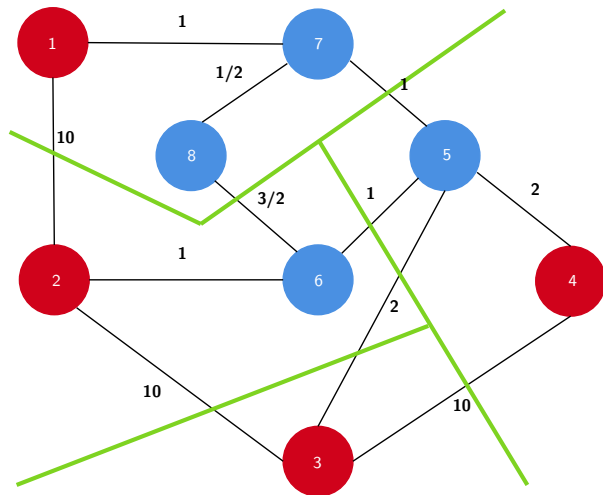
1. compute metric closure $\rightarrow G_1 := (R, E_1, d_1)$
2. compute minimum spanning tree $\rightarrow G_2 := MST(G_1)$
3. replace edges in G_2 by their shortest paths in G $\rightarrow G_3$
4. compute minimum spanning tree $\rightarrow G_4 := MST(G_3)$
5. delete unnecessary edges $\rightarrow T$

Mehlhorn Algorithm

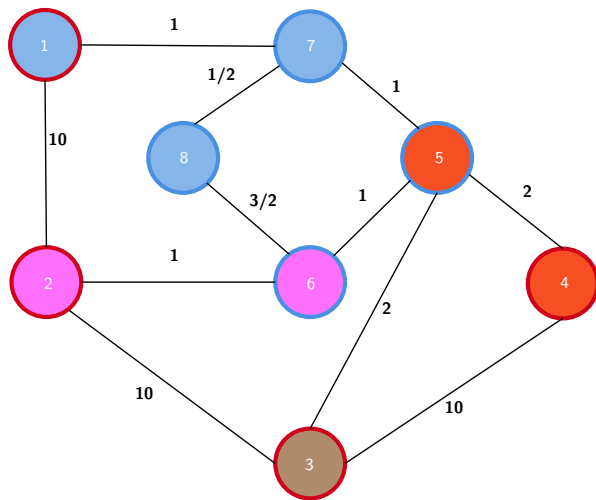
Input graph



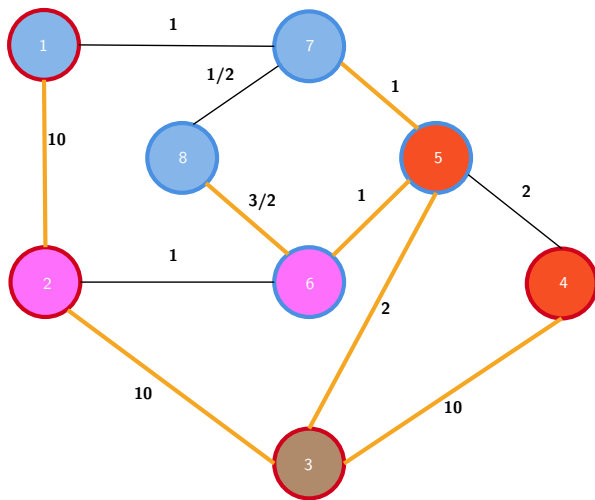
Partition $V \rightarrow$ Voronoi regions



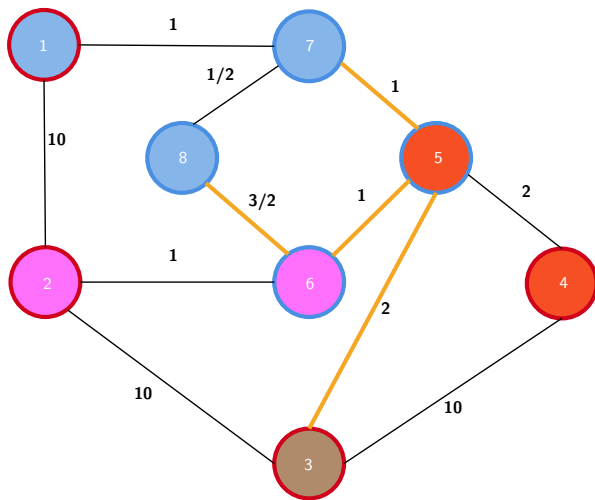
Partition $V \rightarrow$ Voronoi regions



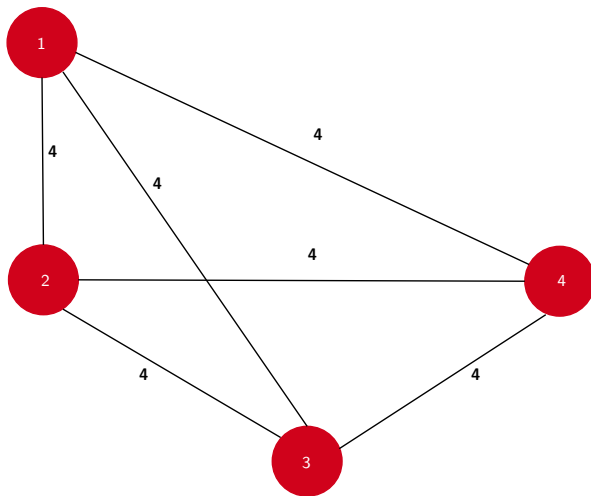
Consider all interregional edges



Choose $|R|$ interregional edges that lead to the shortest interterminal paths



Create a new graph G' induced by terminals



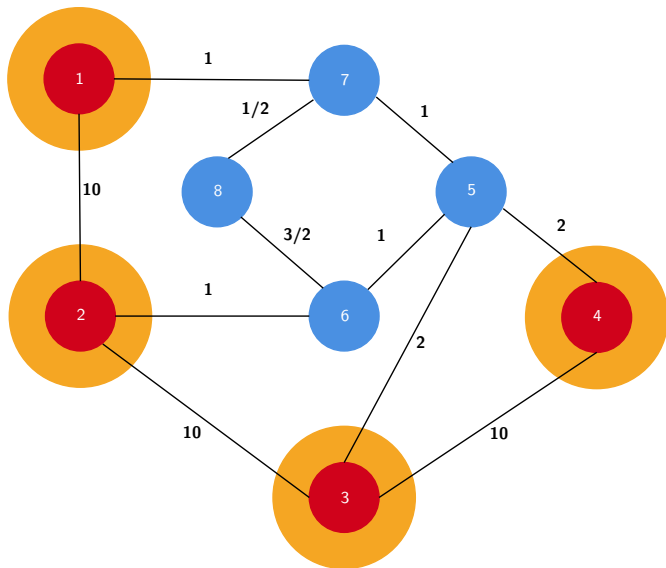
Proceed with steps 2-5 of Kou-Markowsky-Berman Algorithm

- ▶ compute minimum spanning tree $\rightarrow G_2 := MST(G')$
- ▶ replace edges in G_2 by their shortest paths in G $\rightarrow G_3$
- ▶ compute minimum spanning tree $\rightarrow G_4 := MST(G_3)$
- ▶ delete unnecessary edges $\rightarrow T$

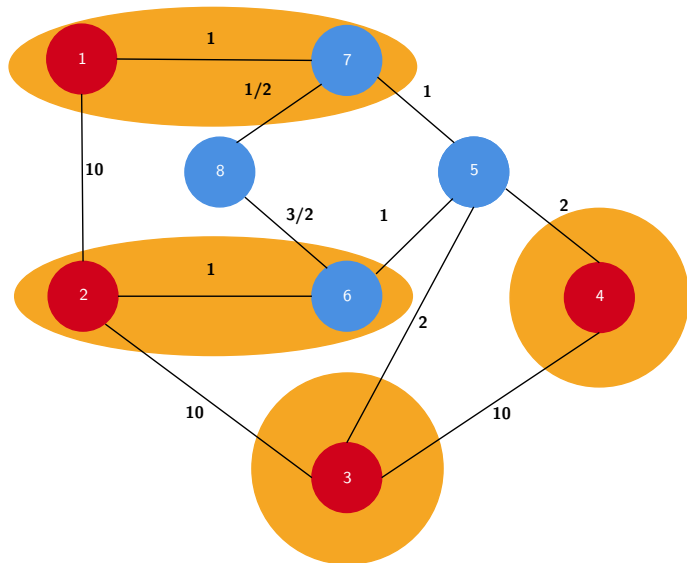
1. construct G' as before $\rightarrow G'$
2. compute minimum spanning tree $\rightarrow G_2 := MST(G')$
3. replace edges in G_2 by their shortest paths in G $\rightarrow G_3$
4. compute minimum spanning tree $\rightarrow G_4 := MST(G_3)$
5. delete unnecessary edges $\rightarrow T$

Primal-Dual Algorithm

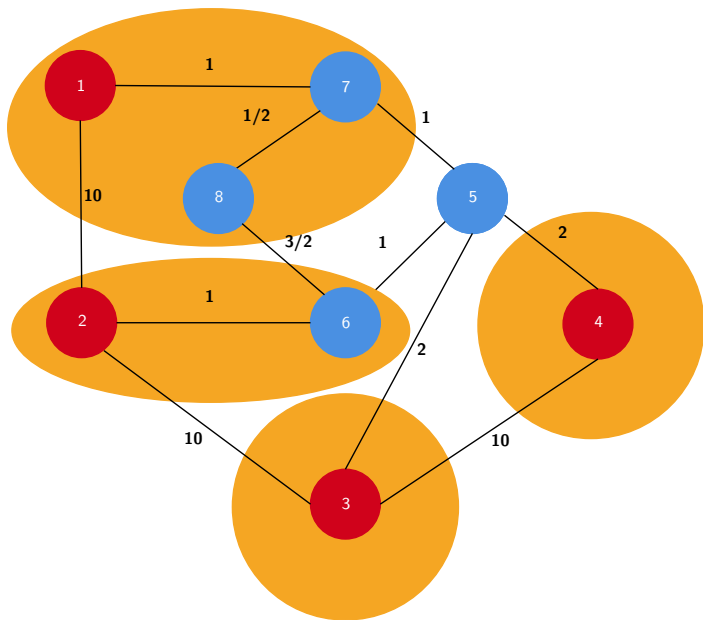
Active sets \mathcal{A}_r for all terminals $r \in R$



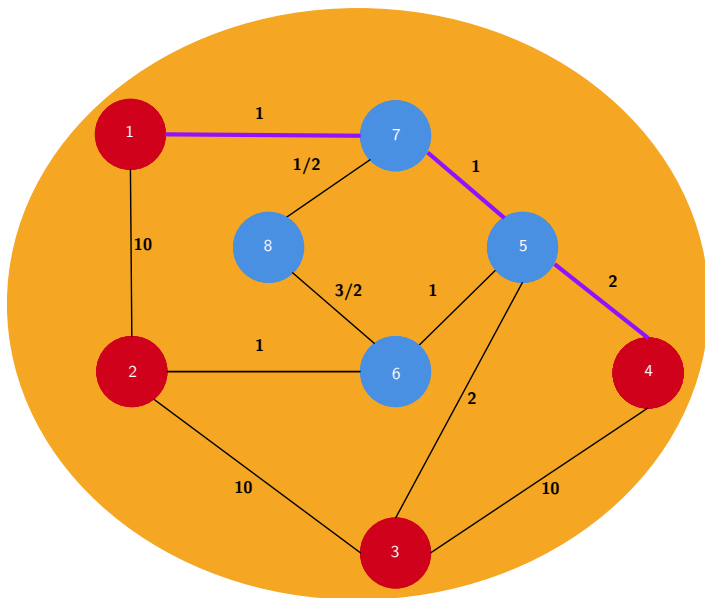
Tight edge: add node to active set



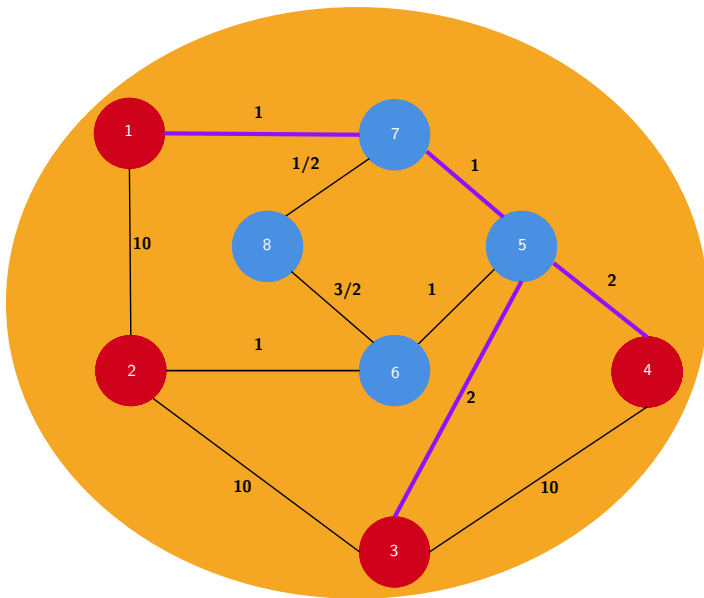
Tight edge: add node to active set



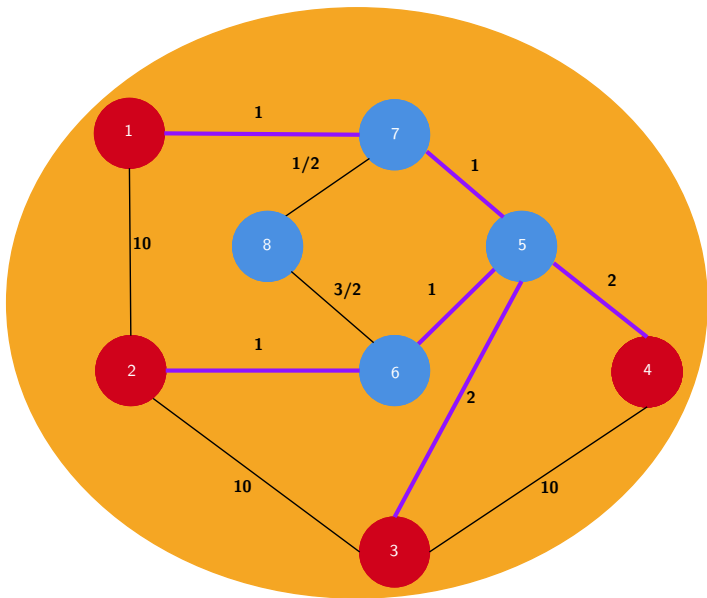
Merge active sets - find connecting path & add it to the primal solution



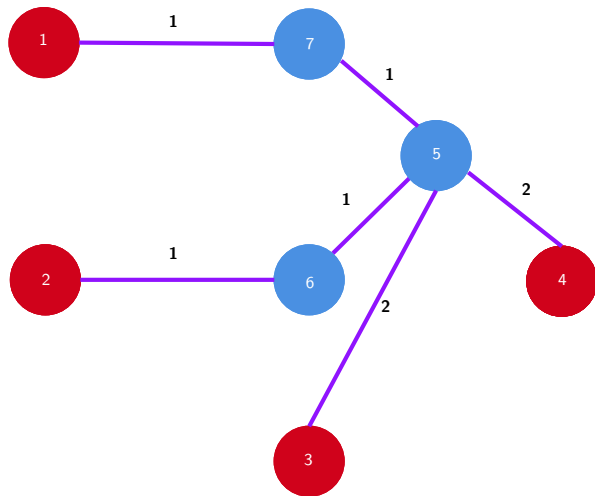
Merge active sets - find connecting path & add it to the primal solution



Merge active sets - find connecting path & add it to the primal solution



Resulting Steiner tree



Deep Reinforcement Learning & Cherrypick

Q-learning:

- ▶ agent learns an action-utility function (Q-function) \rightarrow expected utility of taking an action in a given state: **$Q(\text{state}, \text{action})$**
- ▶ create & update a Q-table (**fixed size!**)

Q-learning:

- ▶ agent learns an action-utility function (Q-function) \rightarrow expected utility of taking an action in a given state: **$Q(\text{state}, \text{action})$**
- ▶ create & update a Q-table (**fixed size!**)

Steiner Tree setting:

- ▶ **state** := nodes included in the current solution
- ▶ **action** := add an uncovered node to the state

Q-learning:

- ▶ agent learns an action-utility function (Q-function) \rightarrow expected utility of taking an action in a given state: **Q(state, action)**
- ▶ create & update a Q-table (**fixed size!**)

Steiner Tree setting:

- ▶ **state** := nodes included in the current solution
- ▶ **action** := add an uncovered node to the state

\implies large # possibilities (**state, action**)

\implies creating & maintaining a Q-table is extremely difficult

Solution: **deep reinforcement learning**:

- ▶ use a deep neural network (NN) to approximate the Q-value function
- ▶ does **not** compute all **Q(state, action)**
- ▶ computes **Q(action)**, i.e. **Q(node)**

- ▶ iteratively grows a connected component C by adding to it a vertex which:
 - ▶ is connected to C in G
 - ▶ has the maximum Q-value

- ▶ iteratively grows a connected component C by adding to it a vertex which:
 - ▶ is connected to C in G
 - ▶ has the maximum Q-value
- ▶ two main parts:
 1. **Graph Embedding Model** \implies 'manipulable' graph format
 2. **Deep Q Network** \rightarrow learn Q-values

- ▶ **similarity matrix** → relationship btw. each node and the 'special' k terminals

- ▶ **similarity matrix** → relationship btw. each node and the 'special' k terminals
- ▶ **message passing** → obtain **vertex embeddings** for all $v \in V$ considering:
 - ▶ similarity matrix
 - ▶ neighbors' embeddings
 - ▶ $v \in R$?
 - ▶ v in the current solution?
- ▶ use vertex embeddings to compute **Q-values** for all $v \in V$

- ▶ **similarity matrix** → relationship btw. each node and the 'special' k terminals
- ▶ **message passing** → obtain **vertex embeddings** for all $v \in V$ considering:
 - ▶ similarity matrix
 - ▶ neighbors' embeddings
 - ▶ $v \in R$?
 - ▶ v in the current solution?
- ▶ use vertex embeddings to compute **Q-values** for all $v \in V$

The vertex embeddings and the Q-values are computed using learnable weights!

- ▶ learn the Q-values (i.e. optimize the weights) using DQN
- ▶ solve a *standard* reinforcement learning task
- ▶ adding a node v to the current solution gives a reward that depends on:
 - ▶ $v \in R$?
 - ▶ $\text{row}(v)$ in the similarity matrix

How do we compute a Steiner tree?

1. Graph embedding model \implies initial Q-values for all nodes
2. While not all terminals are covered:
 - 2.1 Pick node based on policy & add to the current solution
 - 2.2 **If learning mode:**
 - 2.2.1 Optimize loss function \implies update weights used in the graph embedding model & Q-values
 - 2.2.2 Recompute similarity matrix
 - 2.2.3 Recalculate Q-values
3. Remove non-terminals with degree 1

1. Graph embedding model \implies initial Q-values for all nodes
2. While not all terminals are covered:
 - 2.1 Pick node based on policy & add to the current solution
 - 2.2 **If learning mode:**
 - 2.2.1 Optimize loss function \implies update weights used in the graph embedding model & Q-values
 - 2.2.2 Recompute similarity matrix
 - 2.2.3 Recalculate Q-values
3. Remove non-terminals with degree 1

Special ingredients: 2 DQNs, replay memory (details in the documentation)

1. Approximation ratio
2. Running time
3. Early stopping mechanism (Primal-Dual Algorithm)
4. Transfer learning (Cherrypick)
5. Weight sharing (Cherrypick)
6. Sequential learning (Cherrypick)

1. PACE 2018 Challenge
2. small generated graphs (complete, grid, wheel, ladder)

Approximation Ratio - all algorithms - less than 400 vertices

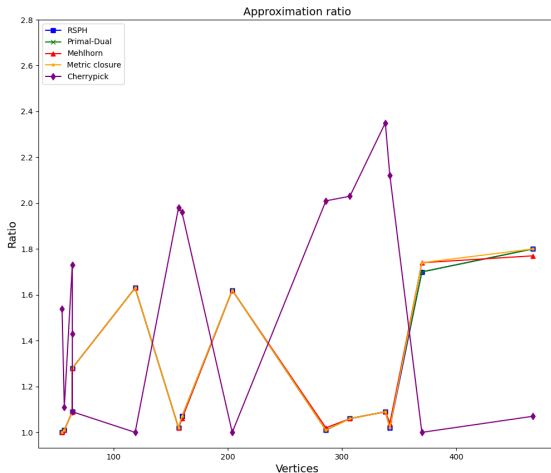


Figure: PACE Challenge Results.

Approximation Ratio - approximation algorithms - less than 7000 vertices

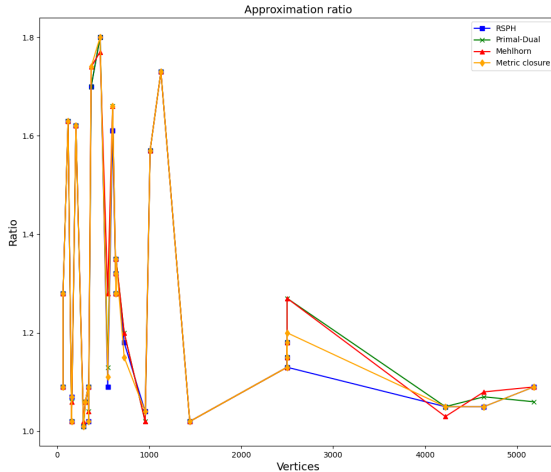


Figure: PACE Challenge Results.

Approximation Ratio - all algorithms - less than 7000 vertices

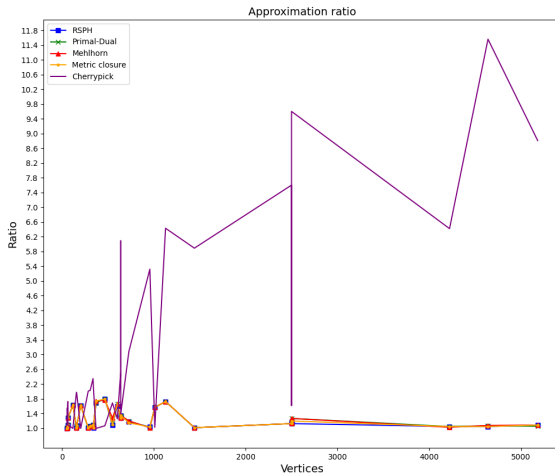


Figure: PACE Challenge Results.

Is it always that bad?

Approximation Ratio - Cherrypick

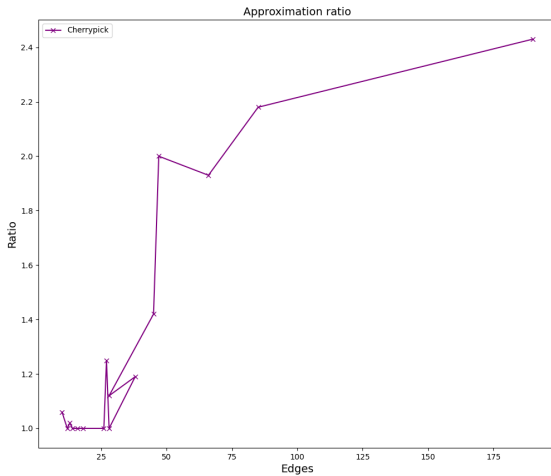


Figure: Results for small graphs.

What would we expect?

Algorithm	Complexity
Mehlhorn	$O(V \log V + E)$
Repetitive Shortest Path	$O(R (E + V \log V))$
KMB (Metric closure)	$O(R (E + V \log V))$
Primal-Dual	$O(V E \log E)$

Table: Complexities of the approximation algorithms

Algorithm	Complexity
Mehlhorn	$O(V \log V + E)$
Repetitive Shortest Path	$O(R (E + V \log V))$
KMB (Metric closure)	$O(R (E + V \log V))$
Primal-Dual	$O(V E \log E)$

Table: Complexities of the approximation algorithms

Mehlhorn < Repetitive Shortest Path = KMB (Metric closure) < Primal-Dual

Running Time - approximation algorithms - less than 400 vertices

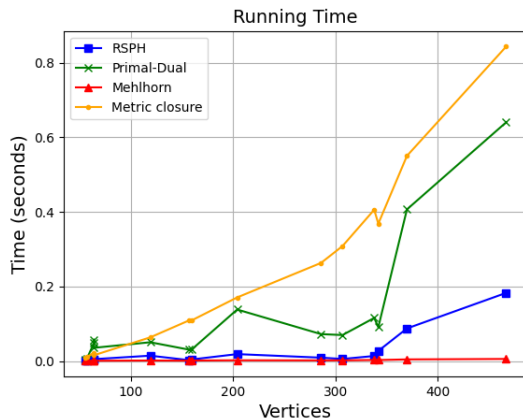


Figure: PACE Challenge Results.

Running Time - approximation algorithms - less than 7000 nodes

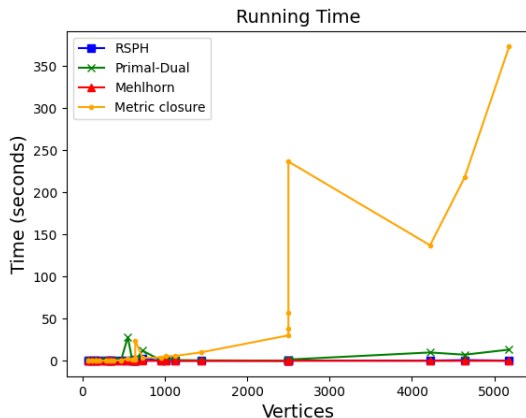


Figure: PACE Challenge Results.

Theory:

Mehlhorn $<$ Repetitive Shortest Path = KMB (Metric closure) $<$ Primal-Dual

Running Time - Approximation Algorithms

Theory:

Mehlhorn $<$ Repetitive Shortest Path = KMB (Metric closure) $<$ Primal-Dual

Practice:

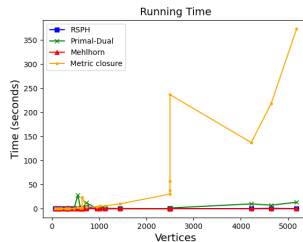
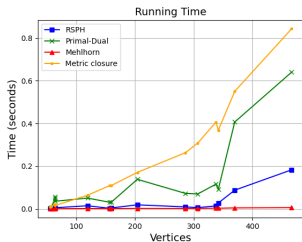


Figure: PACE Challenge Results.

Running Time - Approximation Algorithms

Theory:

Mehlhorn $<$ Repetitive Shortest Path = KMB (Metric closure) $<$ Primal-Dual

Practice:

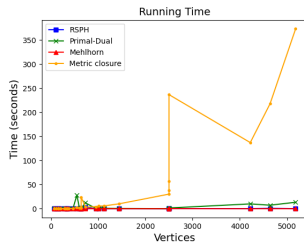
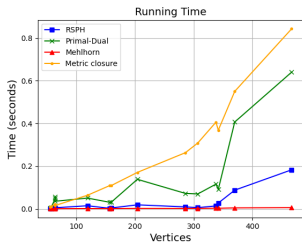


Figure: PACE Challenge Results.

Primal-Dual faster than KMB (Metric closure)?

Running Time - Primal-Dual Algorithm

- ▶ edge sorting - most expensive step
- ▶ sorting algorithm? → Timsort (default in Python)
 - ▶ efficient
 - ▶ takes advantage of any ordering already present in a collection

Running Time - Primal-Dual Algorithm

- ▶ edge sorting - most expensive step
- ▶ sorting algorithm? → Timsort (default in Python)
 - ▶ efficient
 - ▶ takes advantage of any ordering already present in a collection
- ▶ Primal-Dual Algorithm:
 - ▶ subsets of ordered edges are likely to be present in the set of edges
 - ▶ the same 'amount' is added to all dual variables in each iteration

Running Time - Primal-Dual Algorithm

- ▶ edge sorting - most expensive step
- ▶ sorting algorithm? → Timsort (default in Python)
 - ▶ efficient
 - ▶ takes advantage of any ordering already present in a collection
- ▶ Primal-Dual Algorithm:
 - ▶ subsets of ordered edges are likely to be present in the set of edges
 - ▶ the same 'amount' is added to all dual variables in each iteration

→ efficient sorting

Early Stopping Mechanism for Primal-Dual Algorithm

Idea:

- ▶ stop each iteration with probability $p = 0.3$ if at least two active components were already merged
- ▶ make solution feasible:
 - ▶ merge, iteratively, the largest two components $C_1, C_2 \rightarrow$ add shortest path $r_1 \rightarrow r_2$ with $r_i \in C_i, i \in \{1, 2\}$
- ▶ combination Primal-Dual Algorithm & Repetitive Shortest Path Heuristic

Early Stopping Mechanism - Primal-Dual Algorithm - less than 400 vertices

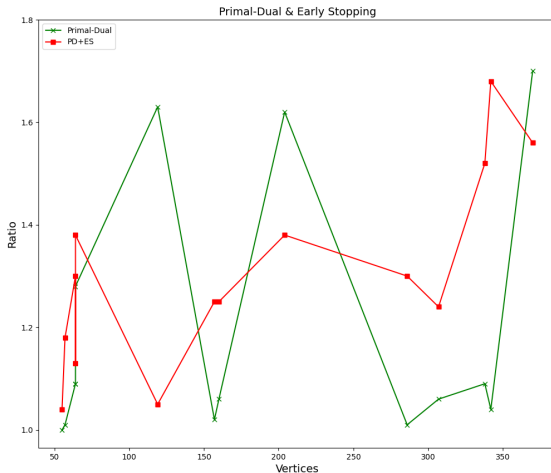


Figure: PACE Challenge Results.

Early Stopping Mechanism - Primal-Dual Algorithm - less than 7000 vertices

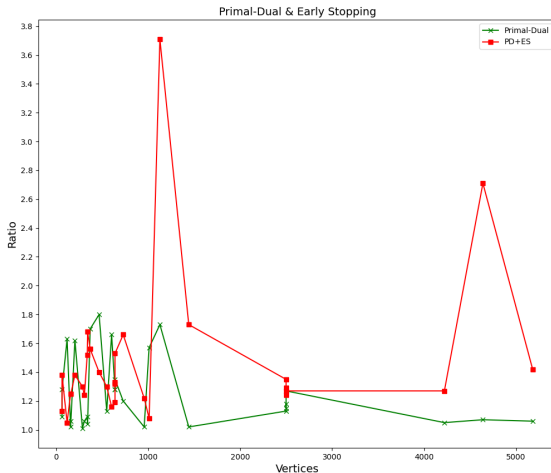


Figure: PACE Challenge Results.

- ▶ good results for small graphs, but no guarantees...
- ▶ drawbacks:
 - ▶ does not take into consideration the 'progress' made between the last iteration in which two components were merged and the iteration that is stopped early
 - ▶ shortest path computation

Cherrypick evaluation scenarios

Consider a graph G and a subgraph $G' \subset G$.

- ▶ randomly initialize weights to solve instance given by $G \rightarrow$ Q-values Q
- ▶ set Q as the initial Q-values when solving instance given by G'
- ▶ apply Cherrypick

Consider a graph G and a subgraph $G' \subset G$.

- ▶ randomly initialize weights to solve instance given by $G \rightarrow$ Q-values Q
- ▶ set Q as the initial Q-values when solving instance given by G'
- ▶ apply Cherrypick

Idea:

- ▶ use learned global knowledge locally

Instance	$ V(G) $	$ V(G') $	$ E(G) $	$ E(G') $	Tr.L.	No Tr.L.
1	17	13	24	16	2.5	2.8
2	17	13	24	13	1	1.6
3	19	14	32	20	1.71	3.23
4	19	14	32	20	3.29	2.52
5	28	28	63	51	2.46	1.82
6	28	20	63	37	2.24	2.71
7	12	8	26	13	1	3
8	12	10	23	15	1	1
9	21	17	27	21	1	1
10	12	10	26	17	1	3

Table: Subgraph results obtained with & without the transfer learning approach.

Idea:

- ▶ improve the **initial solution** computed for an instance in multiple, **independent** algorithm executions
- ▶ 1st run: randomly initialize weights and solve the instance \rightarrow Q-values Q
- ▶ next run: use Q as the initial Q-values and solve the same instance
- ▶ repeat for a fixed number of times (4)

Weight Sharing - Cherrypick

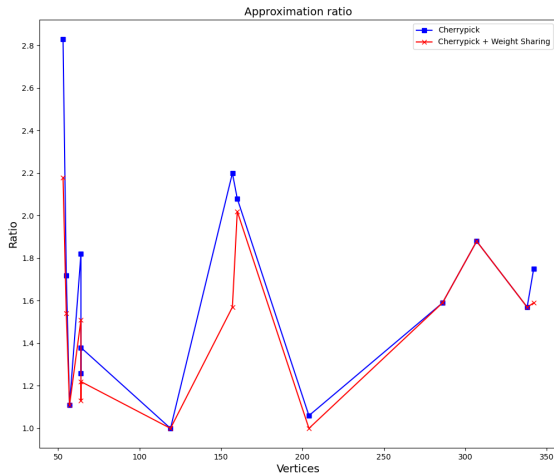


Figure: PACE Challenge Results.

- ▶ Classical approaches achieve similar approximation ratios in practice
- ▶ Mehlhorn is the fastest algorithm in theory & practice
- ▶ KMB (Metric closure) is the slowest in practice (among the classical approaches)
- ▶ Primal-Dual is faster than KMB (Metric closure) in practice due to Timsort

- ▶ Cherrypick approximation ratios grow with the number of edges
- ▶ Primal-Dual - Early stopping mechanism works well for small graphs
- ▶ Cherrypick - Transfer learning approach leads to good results
- ▶ Cherrypick - Weight sharing approach works well, but is computationally expensive